

Quantitative Content Analysis: Lecture 8

Matthias Haber

05 April 2017

Today's outline

- Constructing a document-feature matrix
- Preprocessing
- Regular Expressions

Basic Principles

- Corpus texts are text repositories.
 - Should not have their texts modified as part of preparation or analysis
 - Subsetting or redefining documents is allowable
- A corpus should be capable of holding additional objects that will be associated with the corpus, such as dictionaries, stopword, and phrase lists, etc.
- A document-feature matrix (dfm) is a sparse matrix that is always documents in rows by features in columns
- Encoding of texts should be done in the corpus, and recorded as meta-data in the corpus
 - This encoding should be UTF-8 by default (problem for Windows machines)

Quanteda

quanteda is an R package for managing and analyzing text, created by Kenneth Benoit, Kohei Watanabe, Paul Nulty, Adam Obeng, Haiyan Wang, Ben Lauderdale, and Will Lowe. You can install quanteda from inside RStudio, from the Tools... Install Packages menu, or simply using

```
install.packages("quanteda")
```

You can also install the developers version directly from Github

```
# the devtools package is required  
devtools::install_github("kbenoit/quanteda")
```

Note that on Windows platforms, it is also recommended that you install the RTools suite, and for OS X, that you install XCode from the App Store.

Text analysis workflow

The goal is to simplify text and reduce dimensionality of the dfm created from it. In a nutshell, we want to filter relevant information and discard irrelevant information.

① Creating the corpus

- reading files
- creating a corpus
- adding document variables and metadata

② Defining and delimiting documents

- defining what are “documents” and what are “sentences”

Text analysis workflow (II)

③ Defining and delimiting textual features, using:

- identify instances of defined features (“tokens”) and extract them as vectors
- usually these will consist of terms, but may also consist of:
 - `ngrams` and `skipgrams`, sequences of adjacent or nearby tokens
 - multi-word expressions, through `phrasetoken`
- in this step we also apply rules that will keep or ignore elements, such as
 - punctuation
 - numbers, including or currency-prefixed digits
 - URLs
 - Twitter tags
 - inter-token separators

Text analysis workflow (III)

4 Further feature selection

- Once defined and extracted from the texts (the tokenization step), features may be:
 - removed or kept through use of predefined lists or patterns
 - collapsed by:
 - stemming
 - converting to lower case

Analysis of documents and features

- 1 From a corpus.

These steps don't necessarily require the processing steps above:

- `quanteda::kwic`
- `quanteda::textstat_lexdiv`
- `summary`

Analyzing a dfm

- 2 From a dfm – after dfm on the processed document and features.

dfm	print
convert	removeFeatures
docfreq	similarity
docnames	dfm_sort
featnames	textmodel
textstat_lexdiv	topfeatures
ndoc	dfm_trim
ntoken	weight
plot	settings
show	

Creating and exploring a corpus

The `quanteda` package comes with a built-in set of inaugural addresses from US Presidents. The `summary` command will output the name of each text along with the number of types, tokens and sentences contained in the text.

```
library(quanteda)
summary(data_corpus_inaugural[1:3])
##           Text Types Tokens Sentences
## 1 1789-Washington    626   1540         23
## 2 1793-Washington    96    147          4
## 3   1797-Adams    826   2584         37
oneText <- data_corpus_inaugural[1]
nchar(oneText)
## 1789-Washington
##           8618
```

Tokenization

To tokenize a text is to split it into units, most commonly words, which can be counted and to form the basis of a quantitative analysis. The quantda function `tokenize` can be used on a character vector, a vector of character vectors, or a corpus.

```
tokens <- tokenize("Do not use semicolons. They are  
transvestite hermaphrodites representing  
absolutely nothing. All they do is show  
you've been to college.")  
  
vec <- c(one = 'This is text one',  
        two = 'This is the second text')
```

```
tokenize(vec)
```

Tokenization (II)

The `tokenize` function has a number useful parameters. To remove punctuation, set the `removePunct` argument to be `TRUE`. We can combine this with the `toLower` function to get a cleaned and tokenized version of our text.

```
tokenize(toLower(vec), removePunct = T)
## tokenizedTexts from 2 documents.
## one :
## [1] "this" "is"   "text" "one"
##
## two :
## [1] "this"  "is"    "the"   "second" "text"
```

Using this function with the inaugural addresses:

```
inaugTokens <- tokenize(toLower(inaugTexts))
```

Creating a dfm

Once each text has been split into words, we can use the `dfm` function to create a matrix of counts of the occurrences of each word in each document. `dfm()` works on a variety of object types, including character vectors, corpus objects, and tokenized text objects.

```
library(dplyr)
inaugDfm <- dfm(inaugTokens) %>%
  trim(min_docfreq = 5, min_count = 10) %>%
  weight(type = 'tfidf')
```

Extract corpus objects

To extract texts from a quantda corpus object we can use the `texts` function.

```
mytexts <- texts(subset(inaugCorpus,  
                        President == "Washington"))
```

Adding metadata

If we are interested in analysing the texts with respect to some other variables, we can create a corpus object to associate the texts with this metadata. For example, we can use the `docvars` option to the `corpus` command to record the party with which each text is associated:

```
dv <- data.frame(Party = c('dem', 'rep', 'rep', 'dem'))
recentCorpus <- corpus(inaugTexts[53:56], docvars = dv)
summary(head(recentCorpus))
```

##		Text	Types	Tokens	Sentences
## 1	1997-Clinton	773	2451	111	
## 2	2001-Bush	622	1810	97	
## 3	2005-Bush	772	2325	100	
## 4	2009-Obama	939	2729	110	

Combining features

We can use this metadata to combine features across documents when creating a dfm:

```
partyDfm <- dfm(recentCorpus, groups = 'Party',  
               ignoredFeatures = (stopwords('english')))  
plot(partyDfm, comparison = TRUE)
```


Investigating character vectors

The fundamental type in which R stores text is the character vector. The most simple case is a character vector of length one. The `nchar` function returns the number of characters in a character vector.

```
s1 <- 'This is my example text'
length(s1)
## [1] 1
nchar(s1)
## [1] 23
```

The `nchar` function is vectorized, meaning that when called on a vector it returns a value for each element of the vector.

```
s2 <- c('This is', 'my example text.', 'So imaginative.')
nchar(s2)
## [1] 7 16 15
```

Task 1

Which were the longest and shortest inaugural addresses speeches?

Task 1: Solution

Which were the longest and shortest inaugural addresses speeches?

```
which.max(nchar(inaugTexts))
```

```
## 1841-Harrison
```

```
##           14
```

```
which.min(nchar(inaugTexts))
```

```
## 1793-Washington
```

```
##           2
```

String extraction

It is not possible to index into a string in R:

```
s1 <- 'This is a very informative example sentences.'  
s1[6:9]  
## [1] NA NA NA NA
```

To extract a substring, instead we use the `substr` function.

```
substr(s1, 6, 9)  
## [1] "is a"
```

String extraction (II)

Often we would like to split character vectors to extract a term of interest. This is possible using the `strsplit` function.

```
s1 <- 'split this string'
strsplit(s1, 'this')
## [[1]]
## [1] "split " " string"
```

Consider the names of the inaugural texts:

```
names(inaugTexts)[1:3]
## [1] "1789-Washington" "1793-Washington" "1797-Adams"
parts <- strsplit(names(inaugTexts), '-')
years <- sapply(parts, function(x) x[1])
pres <- sapply(parts, function(x) x[2])
```

Joining character vectors

The paste function is used to join character vectors together. The way in which the elements are combined depends on the values of the sep and collapse arguments:

```
paste('one', 'two', 'three')
## [1] "one two three"
paste('one', 'two', 'three', sep='_')
## [1] "one_two_three"
paste(years[1:3], pres[1:3], sep='-')
## [1] "1789-Washington" "1793-Washington" "1797-Adams"
paste(years[1:3], pres[1:3], collapse='-')
## [1] "1789 Washington-1793 Washington-1797 Adams"
```

Lower case

toLower and toUpper from the quanteda package change the case of character vectors.

```
s1 <- 'NASA sent a rocket into space.'  
  
quanteda::toLower(s1)  
## [1] "nasa sent a rocket into space."  
  
quanteda::toUpper(s1)  
## [1] "NASA SENT A ROCKET INTO SPACE."
```

Comparing character vectors

Character vectors can be compared using the `==` and `%in%` operators:

```
tolower(s1) == toupper(s1)
## [1] FALSE
'apples' == 'oranges'
## [1] FALSE
'pears' == 'pears'
## [1] TRUE

c1 <- c('apples', 'oranges', 'pears')
'pears' %in% c1
## [1] TRUE

c2 <- c('bananas', 'pears')
c2 %in% c1
## [1] FALSE TRUE
```


Searching and replacing within text

The base functions for searching and replacing within text are `grep` and `gsub`. The `grep` command tests whether a pattern occurs within a string:

```
grep('pear', 'these are oranges')  
## integer(0)  
grep('orange', c('apples', 'oranges', 'pears'))  
## [1] 2  
grep('pears', c('apples', 'oranges', 'pears'))  
## [1] 3
```

The `gsub` command substitutes one pattern for another within a string:

```
gsub('oranges', 'apples', 'these are oranges')  
## [1] "these are apples"
```

String replacement

The `stringr` and `stringi` packages provide more extensive and more organized interfaces for string manipulation.

For an overview of the most frequently used functions, see the vignette: <https://cran.r-project.org/web/packages/stringr/vignettes/stringr.html>.

```
library(stringr)
fruits <- c("one apple", "two pears", "three bananas")
str_replace(fruits, "[aeiou]", "-")
## [1] "-ne apple"      "tw- pears"      "thr-e bananas"
str_replace_all(fruits, "[aeiou]", "-")
## [1] "-n- -ppl-"      "tw- p--rs"      "thr-- b-n-n-s"
str_replace(fruits, "([aeiou])", "\\1\\1")
## [1] "oone apple"      "twoo pears"      "threee bananas"
str_replace(fruits, "[aeiou]", c("1", "2", "3"))
## [1] "1ne apple"      "tw2 pears"      "thr3e bananas"
```

Word boundaries

```
words <- c("These are some words. These are more words.")
str_count(words, boundary("word"))
## [1] 8
str_count(words, boundary("sentence"))
## [1] 2

str_split(words, boundary("sentence"))[[1]]
## [1] "These are some words. " "These are more words."
```

Trim whitespace

```
str_trim("  String with trailing and leading white space\t")  
## [1] "String with trailing and leading white space"
```